

more likely to be truthful when there is censorship of misinformation. Under certain conditions, this policy can ultimately lead to the greater spread of misinformation (part (b)).

While provenance and censorship policies have similar features in that both are effective when done properly (and backfire when not), they possess different comparative advantages. For instance, provenance policies tend to perform better when most of the content on the platform is low reliability. In this case, multiple (independent) fact-checks by social media users typically remove this content more efficiently than just a single regulator can. In contrast, because users might share sensational content that contains misinformation even when there is an effective provenance policy, censorship may be more attractive relative to provenance policies in the presence of misinformation carried by sensationalist articles.

6.3 Performance Targets

Another possible regulation, which has recently been proposed by social media platforms (see [Facebook \(2020\)](#)), is to set *performance targets* for their content. As monitoring and removing misinformation is imperfect, a performance target allows the platform some leeway to have “bad” content on their site while still requiring a degree of accountability.

While in Section 6.2 we considered a policy where the regulator was responsible for removing content, a performance target transfers the burden of content removal to the platform itself. In this subsection, we consider how the platform would remove content when misinformation must be kept below a given threshold. We assume that the platform can discard content (by not recommending it to any user), and in doing so, forfeits any potential engagement this content may have with the users.

As in our previous discussion, we assume that sensational articles tend to be more associated with misinformation (see Section 3.2). Formally, we suppose that the sensationalism of every article is drawn as $\kappa \sim G(\cdot|\nu)$ conditional on the article’s veracity ν , where $G(\cdot|\nu = \mathcal{M}) \succeq_{FOSD} G(\cdot|\nu = \mathcal{T})$. Given this, the regulator sets a performance target of λ , which requires the proportion of misinformation shares (to total shares) on the platform to fall below λ .²⁴ The regulator enforces this performance target by auditing the platform and sampling the content to ensure it meets the required standard. Formally, we assume that the regulator has an auditing technology $\alpha \in (0, 1)$, which represents the probability of detecting that the platform has violated its performance target, and if so, the platform incurs a cost due to regulatory fines.

Our next result establishes how stricter performance targets affect the spread of misinformation:

Proposition 7. *There exists an auditing technology cutoff $\alpha^* < 1$ such that:*

(a) *If $\alpha > \alpha^*$, for a fixed sharing network, a stricter performance target (lower λ) reduces the virality of misinformation;*

²⁴This metric for performance comes from Facebook’s own statements on platform standards: “Regulators could say that internet platforms must publish annual data on the ‘prevalence’ of content that violates their policies, and that companies must make reasonable efforts to ensure that the prevalence of violating content remains below some standard threshold” (from [Facebook \(2020\)](#)), with the definition of prevalence being: “We care most about how often content that violates our standards is actually seen relative to the total amount of times any content is seen on Facebook” (from <https://about.fb.com/news/2019/05/measuring-prevalence/>).